

TRUNG THANH TRAN

AI Researcher & Engineer · Foundation Models · Computer Vision · LLMs & Agentic Systems

Hanoi, Vietnam · tranthanhtrung1990@gmail.com · trungthanhtran.github.io
[LinkedIn](#) · [GitHub](#) · [Google Scholar](#) · [ORCID 0000-0002-4630-1146](#) · [Medium](#)

SUMMARY

AI researcher and engineering leader with 10+ years of experience building production AI systems across computer vision, search, speech, and LLM applications. Currently Head of AI at Pixta Vietnam (PIXTA Inc., Japan), leading systems that review 30K+ media assets daily and improve stock-media search conversion. Co-founder of two AI startups — biometric age verification (ClientScan, Oxford) and voice productivity (TakeNote.ai). Current focus: foundation models, efficient inference, LLM evaluation, and production AI systems for search, media review, and regulated domains.

TECHNICAL SKILLS

Vision & Foundation Models: PyTorch, Vision Transformers (DINOv3), InsightFace, face & age estimation, CLIP & vision-language models

LLMs & Agentic Systems: fine-tuning (LoRA / QLoRA, DPO), RAG, LangGraph, Model Context Protocol (MCP), LLM evaluation & fairness

Search, Serving & Efficiency: learning-to-rank (LightGBM / LambdaRank), vLLM, NVIDIA Triton, ONNX, edge / distributed inference, quantization

Languages: Python, SQL · Vietnamese (native), English (professional)

EXPERIENCE

Head of AI & General Engineering Manager, PixtaStock — Pixta Vietnam (PIXTA Inc., Japan) 2019 – Present

- Lead AI research, innovation, and engineering management for the PixtaStock product — setting technical direction and overseeing delivery across all AI/ML projects (computer vision, NLP, search, image generation).
- Built and operate the automated media-review system (30K+ photos, footage, and illustrations per day at >99% precision, ~70% recall), cutting manual-review cost by ~50% — roughly \$100K–\$150K per year.
- Shipped a learning-to-rank search system (LightGBM / LambdaRank), +5% search conversion, plus an LLM-powered assistant saving ~\$10K/year; earlier, a Transformer auto-annotation model (+5% over prior) and a diversity-aware labeling method published at ICICT 2020.

Co-founder & AI Architect — ClientScan Ltd — Oxford, UK Jan 2023 – Present

- Co-founded a biometric-AI company delivering facial-recognition and age-verification for regulated UK retail, hospitality, and gaming (Challenge-25, gambling self-exclusion); products featured by Biometric Update and Hospitality Tech.
- Led and released FaceAge ClientScan, a deployable face-only age and gender model (DINOv3-ViT-L + CORAL), achieving MAE 3.555 on LAGENDA 84K, outperforming MiVOLO v2's reported 3.650 with face crops only, and Avg MAE 4.884 across eight public age datasets.
- Architected the in-the-wild CCTV age-verification pipeline: face-quality gating, robust frame selection, and ByteTrack-based multi-frame aggregation for robustness to pose, lighting, and occlusion.

Co-founder — TakeNote.ai Jun 2023 – Present

- Co-founded a voice-AI product for meeting productivity: accurate transcription, summarization, and knowledge visualization.

Independent Researcher — LLM Fairness & Cultural Bias — Georgia Institute of Technology Dec 2024 – Present

Part-time research collaboration · with Dr. Tarek Naous

- Co-author of CAMELLIA, a benchmark for cultural bias in LLMs across Asian languages; design evaluation methods to surface and reduce where foundation models mis-handle underrepresented cultures.

AI Researcher — Fatima Fellowship — Colorado School of Mines Aug 2024 – Dec 2025

Fellowship research (part-time) · with Dr. Ismet Dagli

- Developed efficiency and memory-optimization techniques for LLMs, vision-language models, and foundation models for resource-constrained and edge hardware.

AI Engineer / Consultant — Heligate JSC & Pi School of AI — Vietnam / Rome, Italy 2018 – 2020

- Heligate: built the first AI team, OCR document classification for a Japanese government client, and a Rasa conversational system. Pi School: conversational voice assistant for Sorgenia (energy utility), featured in Italian media.

Technical Leader & AI Engineer — Panasonic R&D Center — Hanoi, Vietnam 2015 – 2018

- Face recognition (ArcFace, F1 $\approx$ 0.95, 3D anti-spoofing); end-to-end ASR (DeepSpeech2); x-vector / LDA speaker ID and noise-robust wake-word models for smart speakers.

OPEN MODELS & RESEARCH ARTIFACTS

FaceAge ClientScan — public face-only age and gender model release with PyTorch, ONNX, Safetensors, reproducible LAGENDA 84K benchmarking, and full model card; MAE 3.555, CS@5 75.5%, gender accuracy 97.75%. Released on [Hugging Face](#).

SELECTED PUBLICATIONS

Three IGI Global book chapters on generative AI in education and ethics (2023–2024).

Diversity-aware image labeling (multi-model ensemble + Transformers) — ICICT 2020; earlier peer-reviewed GNSS / satellite-navigation research (2013–2014). Full list on Google Scholar.

Preprints: CAMELLIA — LLM cultural-bias benchmark for Asian languages (with Georgia Tech, 2026); MAGNeto — extractive tag summarization (working paper).

AWARDS & HONORS

- United Nations / Italy Long-Term Fellowship, Politecnico di Torino (2013–2014); Summer Scholarship on GNSS, Institute of PNT of Japan, Tokyo (2014).
- Pi School of AI Scholarship + Co-runner-up, Notartel AI Hackathon, Rome (2018); first group in Asia to demonstrate Galileo-only positioning — featured in Inside GNSS and GPS World (2013).
- Runner-up, ICT Research Contest, HUST (2013); graduated ranked 1/120 (top of class).

EDUCATION

Politecnico di Torino — 2nd-Level Specializing Master, Navigation & Related Applications 2013 – 2014

Hanoi University of Science & Technology (HUST) — Engineer's Degree, Information Technology · Ranked 1/120
2008 – 2013

SERVICE & MENTORING

- Technical book reviewer, Manning Publications (AI, Deep Learning, LLMs, Agentic AI); admissions reviewer, MenaML Winter School 2026 (with Google DeepMind, QCRI, KAUST, Apple, InstaDeep).
- Judge, Vietnam AI Grand Challenge Hackathon (2019); mentor to master's students from Kyoto Institute of Technology & Ritsumeikan University.

REFERENCES

Available upon request — including Ryo Kobari (CTO, PIXTA Inc.; CEO, Pixta Vietnam), Robin Paine (former CTO, London Stock Exchange; CEO & co-founder, ClientScan), Dr. Sébastien Bratières (Director of AI, Pi School / Translated), and Dr. Tarek Naous (Georgia Institute of Technology).